WARSAW UNIVERSITY OF TECHNOLOGY

FACULTY OF MATHEMATICS AND INFORMATION SCIENCE

Real-time fraudulent transactions detection

Big Data Analytics

Salveen Singh Dutt (317298)

Karina Tiurina (335943)

Mikołaj Malec (298828)

Patryk Prusak (305794)

Supervisor

mgr inz. Jakub Abelski

Warsaw 2024

Contents

Introduction

This project aims to plan and implement a financial transaction processing system that identifies suspicious and fraudulent activity in real-time. Given the enormous volume of incoming data, the project will utilize big data technologies and advanced machine learning algorithms for anomaly detection.

GitHub repository: https://github.com/salveendutt/Big-Data-Analytics.

Updates from Milestone 1

The project architecture layout was updated to follow the Lambda architecture better. In particular:

- master data storage was moved to the Batch layer;

- included NoSQL database as a storage for data querying in the Serving layer;

- Hive has been changed to Cassandra as the main storage solution;

- The architecture schema was updated to represent the data flow better.

1 High level description

The main idea is to implement automatic transaction processing so that the transaction is blocked for further manual review when a fraudulent activity occurs. The aim is to reduce the end-users financial losses and enhance the security of online payment, ensuring a safer experience for all customers.

The project's main end-users are financial institutions (we will call them 'Managers') and their customers who are executing the payments. Although both categories can benefit from the solution, in our implementation, we will mainly focus on Managers to limit additional data in storage.

The list below contains the main features that we expect to implement for Managers:

1. Fraudulent transactions are automatically highlighted so that it is easier to identify suspicious activity;

2. The history of transactions is stored and available for later review;

3. A dashboard with statistics of fraudulent activity is available and customizable for better localization of issues (e.g., too large amount, unusual location);

4. Anomaly-detection model is continuously updated so that fraud detection utilizes new historical data and is more accurate on future transactions;

5. Data streaming processing and batch jobs are customizable so that the testing of the model's performance is simplified.

2 Data sources

Due to strict security regulations on personal and financial data, it is challenging to find open-source actual transaction data for model training and streaming. Therefore, available synthetic and anonymized datasets will be used. The table below contains a description of the data sources. Each data source is described in more detail in the dedicated subsections.

Data Source	Content	Volume	Fraud, %	Link
1. Fraudulent Transactions Data	Dataset for predicting fraudulent transactions for a financial company.	6,362,620 rows and 10 columns (493.53 MB)	0.13%	Kaggle
2. Credit Card Fraud	Contains features with transactional context.	1,000,000 transactions (58.9 MB)	8.7%	OpenML
3. Credit Card Transactions Synthetic Data Generation	A collection of synthetic credit card transaction data.	1,785,308 transactions; 5,000 customers; (153.66 MB)	3%	Kaggle
4. Credit Card Fraud Detection	Transactions made by credit cards in September 2013 by European cardholders.	284,807 transactions (150.83 MB)	0.17%	Kaggle

Table 1: Data sources

Data-streaming API was implemented from scratch. The assumption was that it would use the above datasets; with a specified time frame, it would choose a random transaction not used for training and push it for further processing. The probability of a fraudulent transaction will be set manually to some high enough constant value for testing purposes. A detailed description of the implemented streaming API can be found in Chapter 3.

The content of the datasets is exceptionally different, which makes it impossible to combine them into a single dataset. Therefore, each dataset will be treated separately for streaming and ML training.

2.1 Fraudulent Transactions Data (Kaggle)

The dataset contains transactions for a financial company, indicating whether they are fraudulent. Data for the case is available in CSV format, which has 6362620 rows and ten columns. The entire column description is presented in Table 2.

The following transformation should be done to pass the dataset to the ML models:

Column	Content	Type
step	maps a unit of time in the real world. In this case, 1 step is 1 hour. Total steps 744 (30 days simulation)	int
type	type of the transaction. Available values: CASH-IN, CASH-OUT, DEBIT, PAYMENT, and TRANSFER	str
amount	amount of the transaction in local currency	float
nameOrig	customer who started the transaction	str
oldbalanceOrg	initial balance before the transaction	float
newbalanceOrig	new balance after the transaction	float
nameDest	customer who is the recipient of the transaction	str
oldbalanceDest	initial balance recipient before the transaction. Note that there is no information for customers that start with M (Merchants)	float
newbalanceDest	new balance recipient after the transaction. Note that there is no information for customers that start with M (Merchants)	float
isFraud	these are the transactions made by the fraudulent agents inside the simulation. In this specific dataset, the fraudulent behavior of the agents aims to profit by taking control of customers' accounts and trying to empty the funds by transferring them to another account and then cashing out of the system	int
isFlaggedFraud	the business model aims to control massive transfers from one account to another and flags illegal attempts. An illegal attempt in this dataset is an attempt to transfer more than 200.000 in a single transaction	int

Table 2: Columns description. Dataset 1: 'Fraudulent Transactions Data' from Kaggle

1. 'type' column values CASH-IN, CASH-OUT, DEBIT, PAYMENT, and TRANSFER transformed to 1, 2, 3, 4, and 5, respectively;

2. a new attribute 'isMerchant' is calculated. 1 if 'nameDest' starts with 'M', 0 - otherwise.

Generally, the dataset and clean and structured, no other preprocessing, except for the described above is necessary. The potential issue is that it is highly unbalanced (less than 1% of fraud transactions) and contains quite a small number of features (6) for model training.

2.2 Credit Card Fraud (OpenML)

This dataset captures transaction patterns and behaviors that could indicate potential fraud in card transactions. The data comprises several features that reflect the transactional context, such as geographical location, transaction medium, and spending behavior relative to the user's history.

The whole dataset is in the numeric form. Therefore, no additional preprocessing is necessary. The only transformation is to rename the target variable to 'isFraud' to match the format

Column	Content	Type
distance_from_home	This is a numerical feature representing the geographical distance in kilometers between the transaction location and the cardholder's home address.	float
distance_from_last_transaction	This numerical attribute measures the distance in kilometers from the location of the last transaction to the current transaction location.	float
ratio_to_median_purchase_price	A numeric ratio that compares the transaction's price to the median purchase price of the user's transaction history.	float
repeat_retailer	A binary attribute where '1' signifies that the transaction was conducted at a retailer previously used by the cardholder, and '0' indicates a new retailer.	[0, 1]
used_chip	This binary feature indicates whether the transaction was made using a chip (1) or not (0).	[0, 1]
used_pin_number	Another binary feature, where '1' signifies the use of a PIN number for the transaction, and '0' shows no PIN number was used.	[0, 1]
online_order	This attribute identifies whether the purchase was made online ('1') or offline ('0').	[0, 1]
fraud	A binary target variable indicating whether the transaction was fraudulent ('1') or not ('0').	[0, 1]

Table 3: Columns description. Dataset 2: 'Credit_Card_Fraud_' from OpenML

of other datasets. The target feature balance is much better than the previous dataset: 8.7% of fraud. An additional complication is that the feature 'amount' is not available. A ratio to the median purchase of the same customer is provided instead.

2.3 Credit Card Transactions Synthetic Data Generation (Kaggle)

This dataset is a collection of synthetic credit card transaction data. The data is designed to mimic the characteristics of actual credit card transactions while ensuring privacy and compliance with data protection regulations such as the General Data Protection Regulation (GDPR). It contains 1,785,308 transactions for 5000 customers.

Like other datasets, it is quite unbalanced, with 3% of fraudulent transactions. The following preprocessing should be done before passing rows to the ML task:

1. 'entry_mode' column values Contactless, Chip, and Swipe transformed to 1, 2, and 3, respectively;

Column	Content	Type
transaction_id	Random string containing specific transactions id	str
post_ts	Date and time of the transaction	str
customer_id	Specific customer id	str
bin	Bank Identification Number	int
terminal_id	Specific terminal id	int
amt	Transaction amount	float
entry_mode	Mode of the transaction. Possible values are Contactless, Chip, and Swipe.	str
fraud	Target variable containing 1 for the fraudulent transaction and 0 otherwise	int
fraud_scenario	Additional label for the transaction. 97% of the dataset has a value of 0. No specific description for each scenario is provided.	int, [0, 1, 2]
mean_amount	Average transaction amount for a specific customer	float
std_amount	Standard deviation of the transaction amount for a specific customer	float
mean_nb_tx_per_day	Mean number of transactions per day for a specific customer	float
customer_bin	Bank Identification Number of a customer	int

Table 4: Columns description. Dataset 3: 'Credit Card Transactions Synthetic Data Generation' from Kaggle

2. 'amt' renamed to 'amount',

3. 'fraud' renamed to 'isFraud',

4. transaction data itself contains only customer ID. Therefore, an additional step is to find the customer and add related features to the output.

2.4   Credit Card Fraud Detection (Kaggle)

The dataset contains transactions made by credit cards in September 2013 by European cardholders. It includes 284,807 transactions, with only 492 (less than 1%) of fraud ones. Due to confidentiality issues, the dataset contains only numerical input variables resulting from a PCA transformation.

This dataset contains a relatively small amount of data compared to others (about 280k transactions). Therefore, we plan to keep it as an additional dataset in case of any issues with others.

Since the data has already been processed, the only necessary transormation is to rename columns 'Amount' and 'Class' to 'amount', 'isFraud' respectively.

Column	Content	Type
Time	The seconds elapsed between the transaction and the first transaction in the dataset	int
V1 ... V28	The principal components obtained with PCA. The original features and more background information about the data are not provided.	float
Amount	Transaction amount	float
Class	Target variable; 1 for fraudulent transaction and 0 otherwise	int, [0, 1]

Table 5: Columns description. Dataset 4: 'Credit Card Fraud Detection' from Kaggle

3 Data acquisition strategy

Due to the lack of publicly available open streaming APIs, we designed and implemented a custom streaming API for real-time fraud detection. Stream API is connected to the NiFi for further data collection and preprocessing. The following technological stack is used for the data acquisition:

1. Python,

2. Flask - for the stream API,

3. Apache NiFi - for the data collection,

4. Docker - for deployment,

5. Apache Kafka - for the data streaming,

6. Kafdrop - for Kafka monitoring,

7. Apache Cassandra - as a main storage solution

8. Zookeeper.

When the server is started, stream API is available on localhost:5000/data/:dataset_id. Data frequency is configurable on the NiFi side and, at the moment, is set to 1 row per 2 seconds for each dataset. The format of each incoming transaction is JSON, containing attributes as described for each dataset in Chapter 2. The streaming API randomly (10%) returns a 500 error to simulate real-world conditions. Example screenshots of the data stream and NiFi are provided in the SKaMP_Tests.pdf file.

Below, we present the NiFi flows for the data acquisition. Processor failures are retried 10 times, and errors are logged with the LogMessage processor with a level of 'error.' The first flow is the overview of the whole process. The three processor groups are responsible for fetching and processing the data. Inside each of the processor groups, there are:

- InvokeHTTP - to fetch the data from the stream API,

- EvaluateJsonPath - to extract the data from the JSON format,

- ReplaceText - to construct a CQL query for the Cassandra,

- PublishKafkaRecord - to send the data to Kafka,
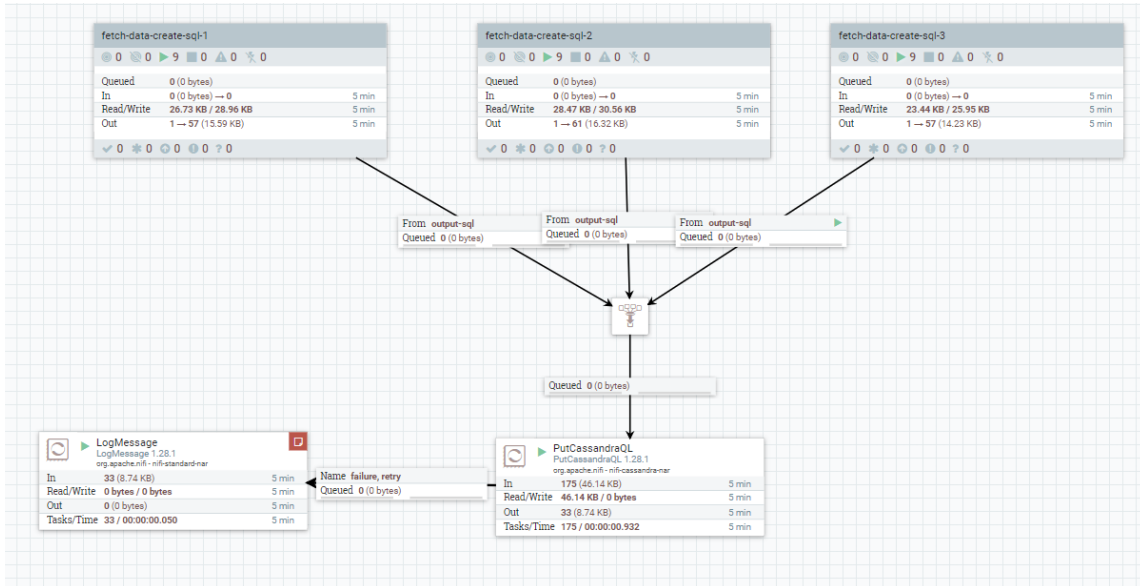
- LogMessage - to log the errors.
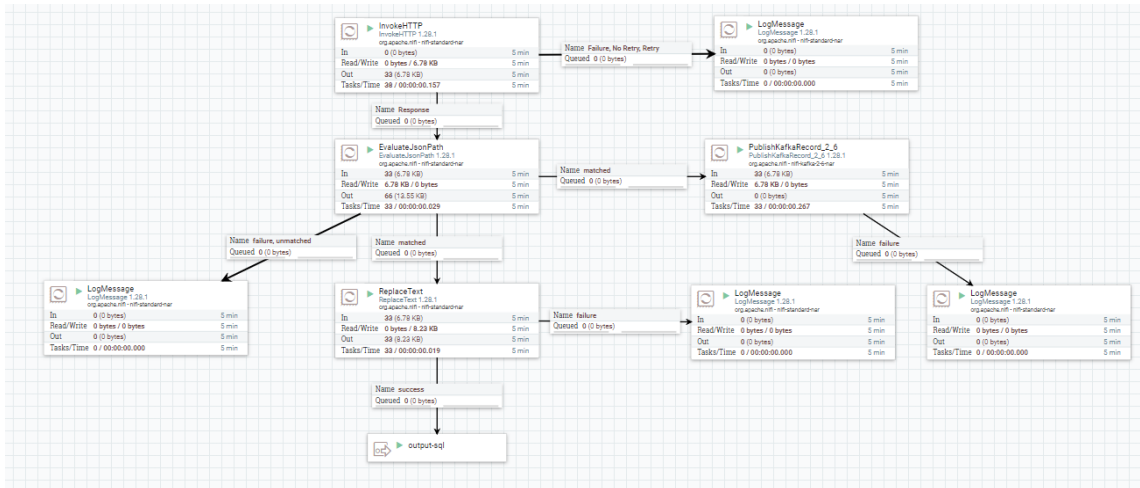

Figure 1: NiFi flow for the data acquisition

Figure 2: Data processing flow in NiFi

4   Data storage strategy

Our input data consists of structured data; there is no need to store particular data types, such as images, text files, audio, etc. Therefore, we decided to utilize an SQL-like (CQL) data warehouse system that enables analytics at a massive scale - Apache Cassandra.

As described in Chapter 2, our datasets contain entirely different sets of features, which makes it impossible to combine them in a single table. All incoming transactions will be divided into four groups for each dataset. The code block below represents an example of table creation for dataset 1 of the project.

Listing 1: Apache Cassandra table creation

```
CREATE TABLE dataset1 (
    step INT,
    type TEXT,
    amount DECIMAL,
    isFlaggedFraud INT,
    isFraud INT,
    nameDest TEXT,
    nameOrig TEXT,
    newbalanceDest DECIMAL,
    newbalanceOrig DECIMAL,
    oldbalanceDest DECIMAL,
    oldbalanceOrg DECIMAL,
    PRIMARY KEY (step, nameOrig)
);
```

The preliminary list of all tables in the storage is the following:

1. dataset1 - raw data for the dataset 1;

2. dataset1_processed - processed data for the dataset 1;

3. dataset2 - raw data for the dataset 2;

4. dataset2_processed - processed data for the dataset 2;

5. dataset3 - raw data for the dataset 3;

6. dataset3_processed - processed data for the dataset 3;

7. dataset4 - raw data for the dataset 4;

8. dataset4_processed - processed data for the dataset 4.

Additionally, we will use MongoDB to store prepared views for fast querying during the data presentation. Typically, it replicates the processed data from batch and streaming layers. This includes predicted labels for new transactions and an indication of whether data was used for training.

5   Project architecture

We are planning to implement the project based on Lambda Architecture. The main data processing will be divided into three layers:

1. Speed Layer (streaming)

   - Data preprocessing including transformation to a specific format;
   - Real-time fraud detection on all of the incoming transactions.

2. Batch Layer

   - Data processing and filtering for the model training
   - ML model training with a fixed schedule (e.g., every 10 minutes)

3. Serving Layer

   - Stores processed real-time and batch data in NoSQL for fast querying
   - Client interface highlighting fraud transactions, accepting/blocking transactions;
   - Data visualization with customizable filters

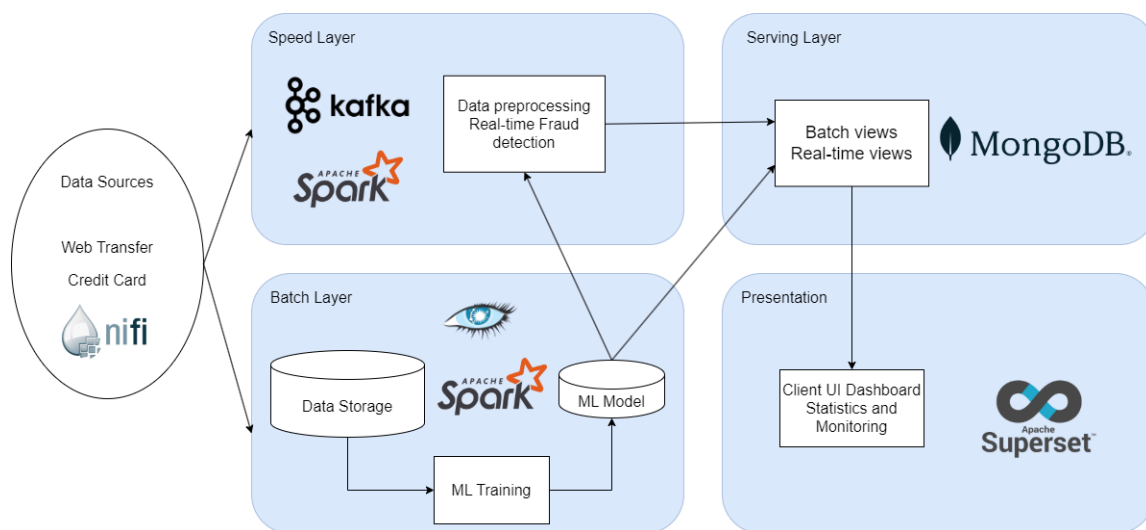Figure 1 shows an outline of the project architecture.

Figure 3: Project architecture

The following Big Data platforms will be used:

- Apache NiFi: to collect and distribute the data from different sources;

- Apache Cassandra as a data storage;

- Apache Kafka: to work with the streaming data;

- Apache Spark: to make the batch processing and model training;

- MongoDB: to store prepared batch views and real-time views for fast querying;

- Apache Superset (to be agreed with the supervisor): for data analysis on the user interface. If the service is not approved, UI will be implemented from scratch using a JS framework, e.g., React.

6 Planned ML Tasks

This section explores potential machine learning tasks that could be applied to the available fraud and anomaly detection datasets. The primary focus will be identifying fraudulent transactions, developing anomaly detection models, and leveraging data-driven insights for effective fraud management.

6.1 Fraud Detection

Fraud detection is a key use case for the provided datasets, where the goal is to classify transactions as fraudulent or legitimate. This can be achieved using various supervised learning techniques. Fraudulent transactions are typically rare, so special care must be taken to handle the class imbalance. Planned steps include:

- **Model Selection:** Applying binary classification models such as Logistic Regression, Random Forest, Gradient Boosted Trees (e.g., XGBoost, LightGBM), and Neural Networks. These models will be trained to predict the 'isFraud' label in the datasets.

- **Evaluation Metrics:** Since fraud detection deals with imbalanced data, standard metrics such as accuracy are insufficient. Precision, recall, F1-score, and the area under the Precision-Recall curve (AUC-PR) will be used to evaluate model performance.

- **Imbalanced Data Handling:** Techniques such as SMOTE (Synthetic Minority Oversampling Technique), undersampling, or cost-sensitive learning will be explored to address class imbalance.

6.2 Anomaly Detection

Anomaly detection is complementary to supervised fraud detection, focusing on identifying unusual transactions that deviate significantly from normal behavior. This approach is beneficial for uncovering previously unseen fraud patterns. Proposed methods include:

- **Unsupervised Learning:** Techniques like Isolation Forests, k-means clustering, and DBSCAN will be applied to identify anomalous transactions based on transaction features.

- **Autoencoders:** Neural network-based autoencoders will be used to reconstruct transaction data, with high reconstruction errors indicating potential anomalies.

- **Density Estimation:** Probabilistic models such as Gaussian Mixture Models (GMM) will be employed to estimate the likelihood of transactions and flag those with low likelihoods as anomalies.

- **Validation:** Anomalies detected by these methods will be cross-validated against the 'isFraud' label to evaluate their effectiveness.

7   Planned batch and stream processing

This chapter describes the planned implementation of batch and stream processing flows in detail as part of the Lambda Architecture outlined in the previous section.

7.1   Batch Processing Workflow

Batch processing will focus on historical data analysis, feature engineering, and machine learning model training. The workflow for batch processing is as follows:

1. Data Ingestion:

   - Periodic ingestion of historical transactions from data storage sources (Apache Cassandra).
   - Data preprocessing to validate and filter incomplete or invalid entries.
   - Aggregation of transactional data for customer-level or terminal-level insights.

2. Data Transformation:

   - Transformation of raw transaction data into structured formats suitable for ML model training.
   - Feature engineering, such as creating rolling statistics (e.g., transaction frequency, average amount per customer, or terminal-specific patterns).

3. Model Training:

   - Training ML models using the preprocessed data with a fixed schedule (e.g., every 30 minutes or daily, depending on model requirements).
   - Storing the trained models in a model repository for use in the serving and streaming layers.

7.2   Stream Processing Workflow

The stream processing layer will handle real-time fraud detection and serve as the system's speed layer. The workflow for stream processing is as follows:

1. Data Ingestion:

- Continuous ingestion of live transaction data via Apache Kafka.
- Real-time preprocessing, including feature transformation (e.g., converting categorical variables, handling missing values).

2. Real-Time Fraud Detection:

- Use of trained ML models to perform predictions on incoming transactions.
- Updating fraud detection logic dynamically as new models are trained in the batch layer.

3. Event Streaming:

- Streaming results to MongoDB or another serving layer for downstream consumption.
- Flagging suspicious transactions for further analysis.

7.3   Serving Layer Workflow

The serving layer will provide a unified view of real-time and batch-processed data and offer tools for decision-making. The workflow is as follows:

1. Unified Data Storage:

- MongoDB will be the primary storage for batch and real-time data, enabling fast querying and retrieval.

2. Client Interfaces:

- Interfaces for visualizing flagged transactions in real-time.
- Fraud management tools to allow users to accept or block flagged transactions.

3. Data Visualization:

- Customizable dashboards for exploring transaction patterns, fraud rates, and model performance.
- Apache Superset will be evaluated as the primary visualization tool; if not approved, a custom solution will be implemented using React.

7.4   Integration of Batch and Stream Processing

The Lambda Architecture ensures seamless integration between the batch and stream processing layers:

- ML models trained in the batch layer will be deployed to the stream processing layer for real-time fraud detection.

- Features engineered in the batch layer can inform the design of streaming features to ensure consistency.

- The serving layer will present a unified view of historical and real-time data, enabling holistic decision-making.

8 Planned way of presenting the results

In this architecture, we propose to use Apache Superset as the primary tool for presenting analytical results and monitoring system performance. Superset is an open-source data visualization and business intelligence platform well-suited for modern data pipelines due to its extensive support for a wide range of data sources and ease of integration into existing architectures.

8.1 Integration of Apache Superset with the Current Architecture

As illustrated in the architecture diagram:

- The **Batch Layer**, powered by Apache Cassandra, processes large-scale data and aggregates it into queryable formats.

- The **Speed Layer**, which uses Kafka for real-time data ingestion and Spark for processing, produces real-time views stored in MongoDB. Superset can connect to MongoDB to display these real-time analytics, provided an appropriate SQL-like connector (e.g., MongoDB BI Connector) is configured.

- The **Serving Layer** facilitates the querying of both batch and real-time views, making this data accessible to Superset for visualization in unified dashboards.

By integrating Superset with batch and speed layers, the system can provide users with comprehensive dashboards that include historical trends and real-time updates. This combination is particularly beneficial for monitoring critical use cases such as fraud detection, as it allows stakeholders to view both immediate alerts and long-term patterns.

8.2 Advantages of Using Apache Superset

- **Ease of Use**: Superset provides an intuitive drag-and-drop interface for creating visualizations and dashboards, making it accessible to non-technical users.

- **Real-Time and Batch Data Integration**: Combining data from the speed and batch layers into unified dashboards enables comprehensive analytics and monitoring.

- **Custom Visualization**: Superset supports a variety of visualizations, allowing users to explore data in formats that best suit their needs, from basic bar charts to advanced geographic maps.

- **Scalability**: Being lightweight and web-based, Superset can handle a growing volume of data as the architecture scales.

- **Open Source and Extensibility**: Superset can be customized and extended to fit specific requirements as an open-source platform.

8.3 Disadvantages and Challenges

While Superset provides significant benefits, there are a few challenges and limitations to consider:

- **Indirect Support for Kafka**: Superset does not natively connect to Kafka. Before it can be visualized, real-time data from Kafka must first be stored in a queryable data store (e.g., MongoDB or a similar OLAP engine).

- **MongoDB Support**: Superset does not have native support for MongoDB. An SQL-like connector (such as the MongoDB BI Connector) or data transformation into another supported format is required, adding some complexity.

- **Performance Considerations**: For real-time monitoring, performance can be a bottleneck if the underlying data stores (e.g., MongoDB or Cassandra) are not optimized for frequent querying by Superset.

- **Learning Curve for Advanced Features**: While the primary interface is user-friendly, advanced configurations (e.g., complex filters and custom SQL queries) may require technical expertise.

9 Tasks assignment

The table below contains the list of team members and the preliminary allocation of tasks to team members.

Team member	Tasks	Supporter
Salveen Singh Dutt	Batch processing of the historical data for up-to-date model training (Batch Layer).	Karina Tiurina
Karina Tiurina	Fraud detection model training and fine-tuning; Data stream processing (Speed Layer).	Salveen Singh Dutt
Patryk Prusak	Data ingestion, collection, and preprocessing.	Karina Tiurina, Salveen Singh Dutt
Patryk Prusak	Data visualization and configuration on the UI (Serving Layer).	Karina Tiurina, Salveen Singh Dutt

Table 6: Tasks assignment

References

[1] Chitwan Manchanda, *Fraudulent Transactions Data*, Kaggle, 2022. Available at: `https://www.kaggle.com/datasets/chitwanmanchanda/fraudulent-transactions-data`.

[2] *Credit_Card_Fraud_*, OpenML, 2024. Available at: `https://www.openml.org/search?type=data&status=active&id=45955&sort=runs`.

[3] Carlos José, *Credit Card Transactions Synthetic Data Generation*, Kaggle, 2024. Available at: `https://www.kaggle.com/datasets/cgrodrigues/credit-card-transactions-synthetic-data-generation?select=transactions_df.csv`.

[4] Machine Learning Group - ULB and Andrea, *Credit Card Fraud Detection*, Kaggle, 2018. Available at: `https://www.kaggle.com/datasets/mlg-ulb/creditcardfraud`.

[5] *Lambda Architecture*, Cazton. Available at: `https://cazton.com/consulting/enterprise/lambda-architecture`.

[6] OpenAI, *ChatGPT: Language Model*, Available at: `https://chatgpt.com`.